



Kafka Conversational AI (RAG-Based System)

A Retrieval-Augmented Generation (RAG) conversational system that simulates Franz Kafka's voice and worldview using his literary works, letters, and biographical materials.

Built with **LangChain** + **ChromaDB** + **Ollama** + **Streamlit**.



Project Overview

This project implements a conversational AI system that:

- Retrieves relevant passages from Kafka's works
- Grounds responses in real textual sources
- Simulates Kafka's tone and existential style
- Displays source transparency in the UI

The system is designed for:

- Academic experimentation
 - RAG architecture practice
 - Persona-based LLM systems
 - Transparent source-backed AI conversations
-



Architecture

User Input

→ Streamlit UI
→ `kafka_rag_answer()`
→ ChromaDB Retriever
→ Ollama LLM
→ Grounded Response + Sources
→ UI Display (Answer + Source Panel)

Core Components

1 Ingestion Pipeline

- Load Kafka texts (PDF / TXT)
- Split into chunks
- Add metadata (author, work, type, chunk_id)

- Generate embeddings
- Store in ChromaDB

2 Retrieval Layer

- Semantic similarity search
- Optional metadata filtering
- Returns top-k relevant chunks

3 Generation Layer

- Prompt enforces grounding
- Uses retrieved context only
- Produces stylistically aligned output

4 Streamlit Interface

- Chat interface
- Source transparency panel
- Session state memory
- Example prompts
- Conversation statistics

📁 Project Structure

```
project/
|
├── app.py                  # Streamlit UI
├── testing.py               # RAG logic (kafka_rag_answer)
├── ingest.py                # Text loading & DB creation
├── chroma_db/              # Persistent vector database
├── requirements.txt
└── README.md
```

⚙️ Installation

1 Clone the repository

```
git clone <repo-url>
cd kafka-rag
```

2 Create virtual environment

```
python -m venv lang_env  
lang_env\Scripts\activate
```

3 Install dependencies

```
pip install -r requirements.txt
```

4 Install and run Ollama

Download Ollama and pull a model:

```
ollama pull llama3
```

Make sure Ollama is running.

Running the Application

```
streamlit run app.py
```

Open the browser at:

```
http://localhost:8501
```

Metadata Design

Each chunk stored in ChromaDB includes metadata like:

```
{  
    "author": "Franz Kafka",  
    "work": "Metamorphosis",  
    "type": "novel",
```

```
    "chunk_id": 12  
}
```

This enables:

- Filtering by work
 - Transparent source display
 - Structured retrieval analysis
-

RAG Design Principles

- Responses must be grounded in retrieved passages
 - No hallucinated content outside archive
 - If context is insufficient → acknowledge limitation
 - Persona styling layered *after* grounding
-

Debugging Tips

Check stored works

```
collection.get(include=["metadata"])
```

Check retrieved chunks

```
for doc in docs:  
    print(doc.metadata)  
    print(doc.page_content[:200])
```

If preview shows only .'

- Verify ingestion pipeline
 - Check chunking
 - Ensure text was loaded correctly
-

Future Improvements

- Streaming token-by-token responses

- Similarity score display in UI
 - Confidence score estimation
 - Better metadata filtering logic
 - Hybrid search (BM 2 5 + embeddings)
 - Multi-work comparison mode
-



Educational Value

This project demonstrates:

- Practical RAG implementation
 - Vector database debugging
 - Persona-conditioned generation
 - Retrieval transparency design
 - Streamlit conversational UI engineering
-



Disclaimer

This system simulates Kafka's literary voice for educational purposes.

Responses are generated by a language model and are not authentic historical writings.



Author

Built as an advanced RAG experimentation project.

"I write differently from what I speak, I speak differently from what I think..."